

# Using the OpenVINO™ Toolkit for Deploying Accelerated Deep Learning Applications – Part2 [2021.4]

July 2021

The Intel logo, consisting of the word "intel" in a lowercase, sans-serif font, with a registered trademark symbol (®) to its upper right. The logo is positioned at the bottom left of the slide, partially overlapping a decorative graphic of several blue squares of varying sizes arranged in a stepped pattern.

intel®

# Agenda

## Part 1: OpenVINO Workshop (110mins):

- Demos on DevCloud
  - Post-Training Optimization Tool
  - DL Workbench
  - DL Streamer
- 
- Part2: Q & A(10mins)

# Notices and Disclaimers

- Performance varies by use, configuration and other factors. Learn more at [www.Intel.com/PerformanceIndex](http://www.Intel.com/PerformanceIndex).
- Performance results are based on testing as of dates shown in configurations and may not reflect all publicly available updates. See backup for configuration details. No product or component can be absolutely secure.
- Your costs and results may vary.
- Intel technologies may require enabled hardware, software or service activation.
- All product plans and roadmaps are subject to change without notice.
- Intel disclaims all express and implied warranties, including without limitation, the implied warranties of merchantability, fitness for a particular purpose, and non-infringement, as well as any warranty arising from course of performance, course of dealing, or usage in trade.
- Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.

# Intel® DevCloud for the Edge Demo

<https://software.intel.com/content/www/us/en/develop/tools/devcloud/edge/build/sample-apps.html>

July 2021



# Post-Training Optimization Tool

July 2021

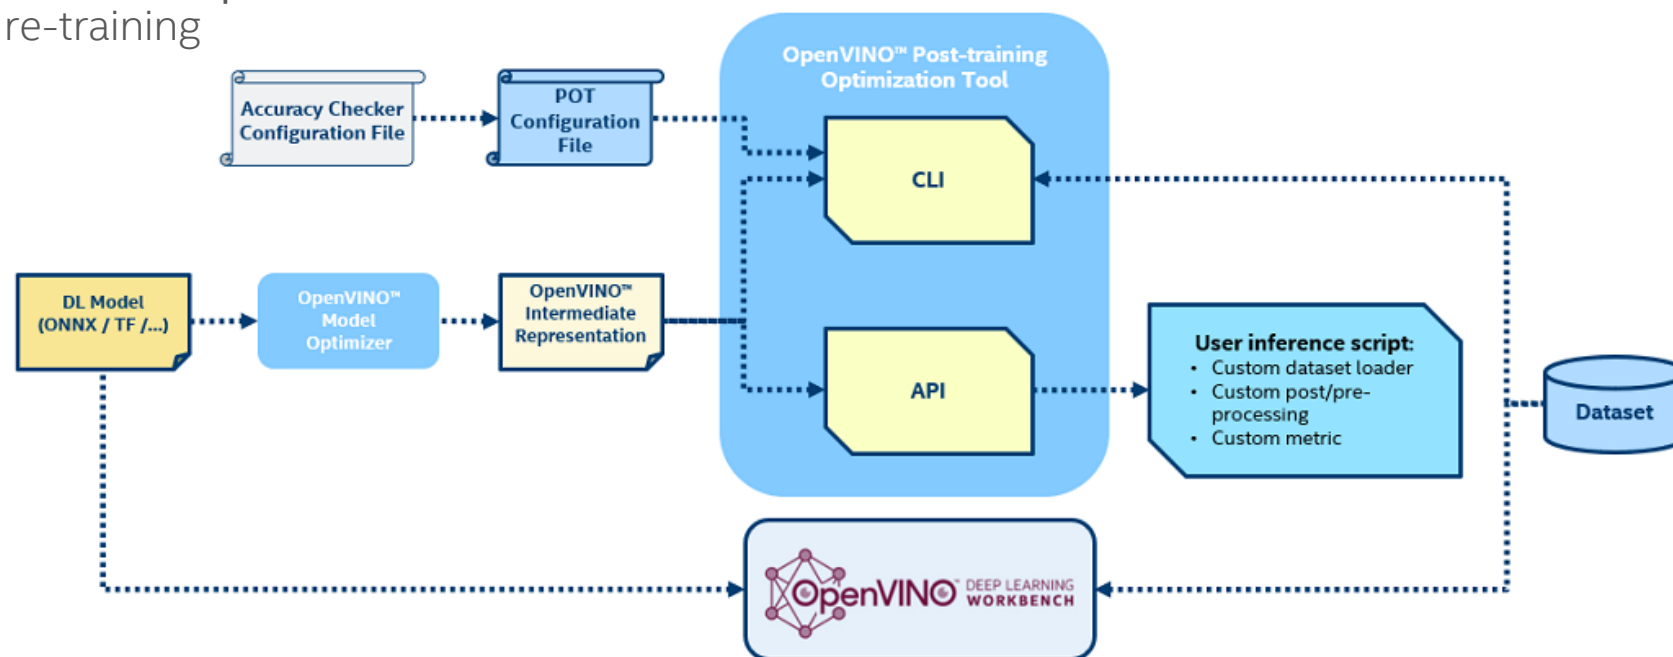


intel<sup>®</sup>

# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

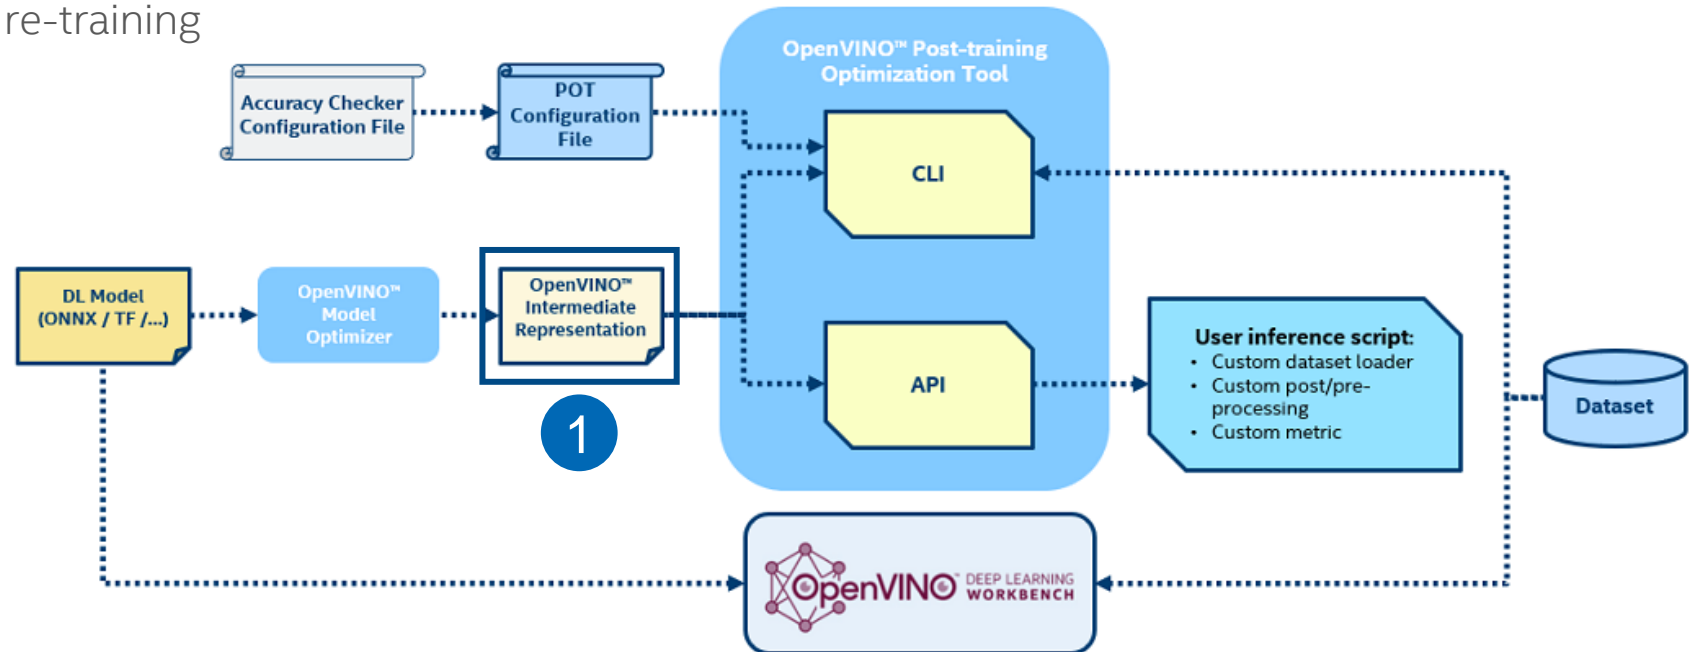
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

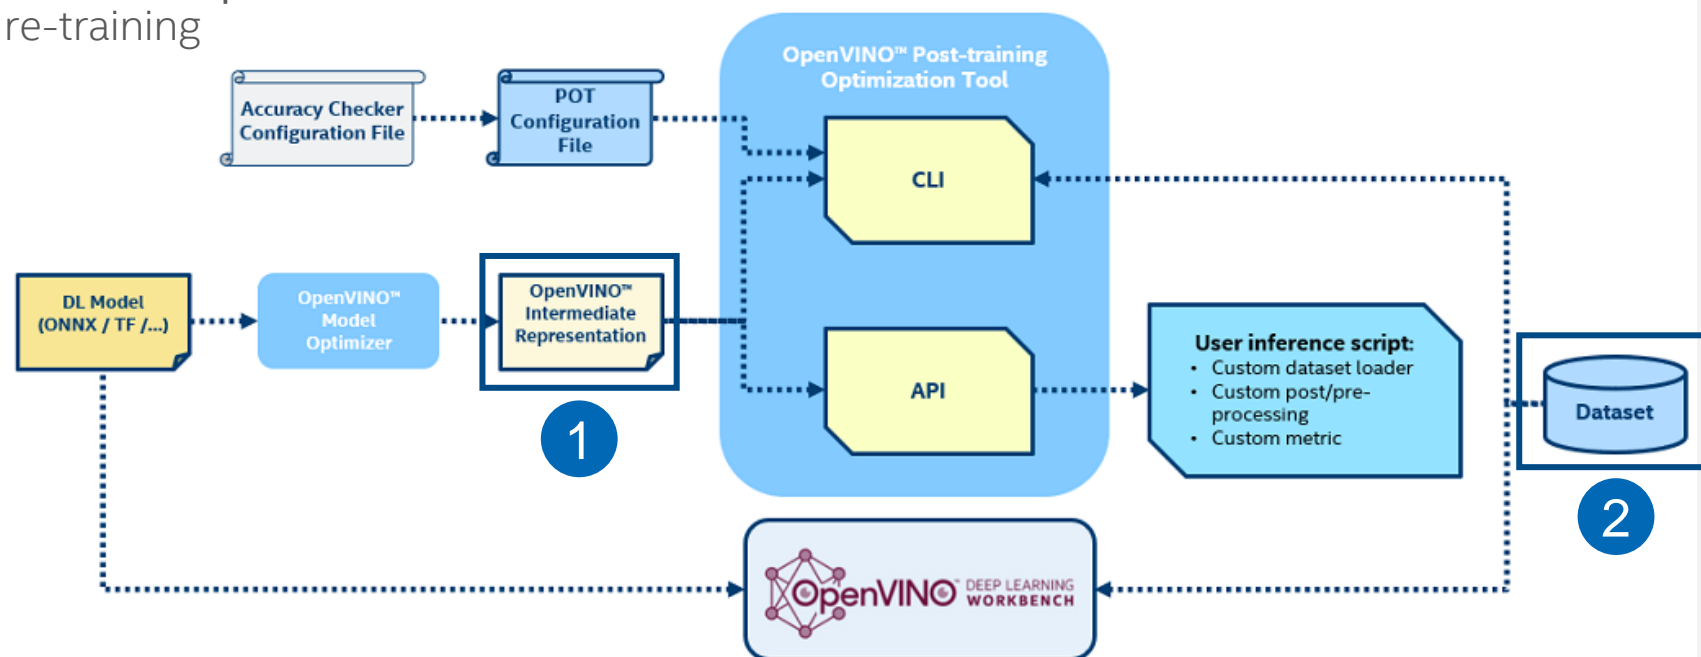
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.

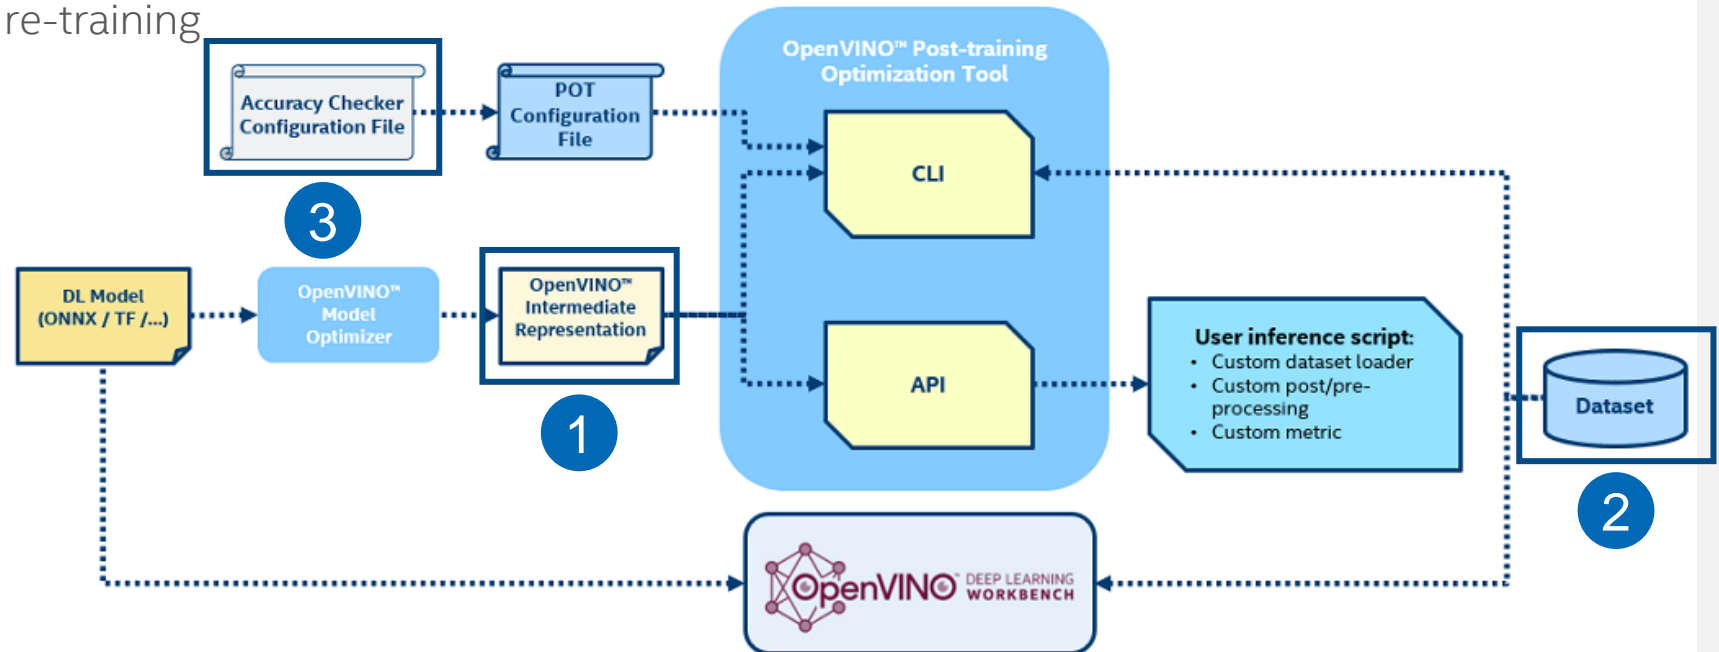




# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

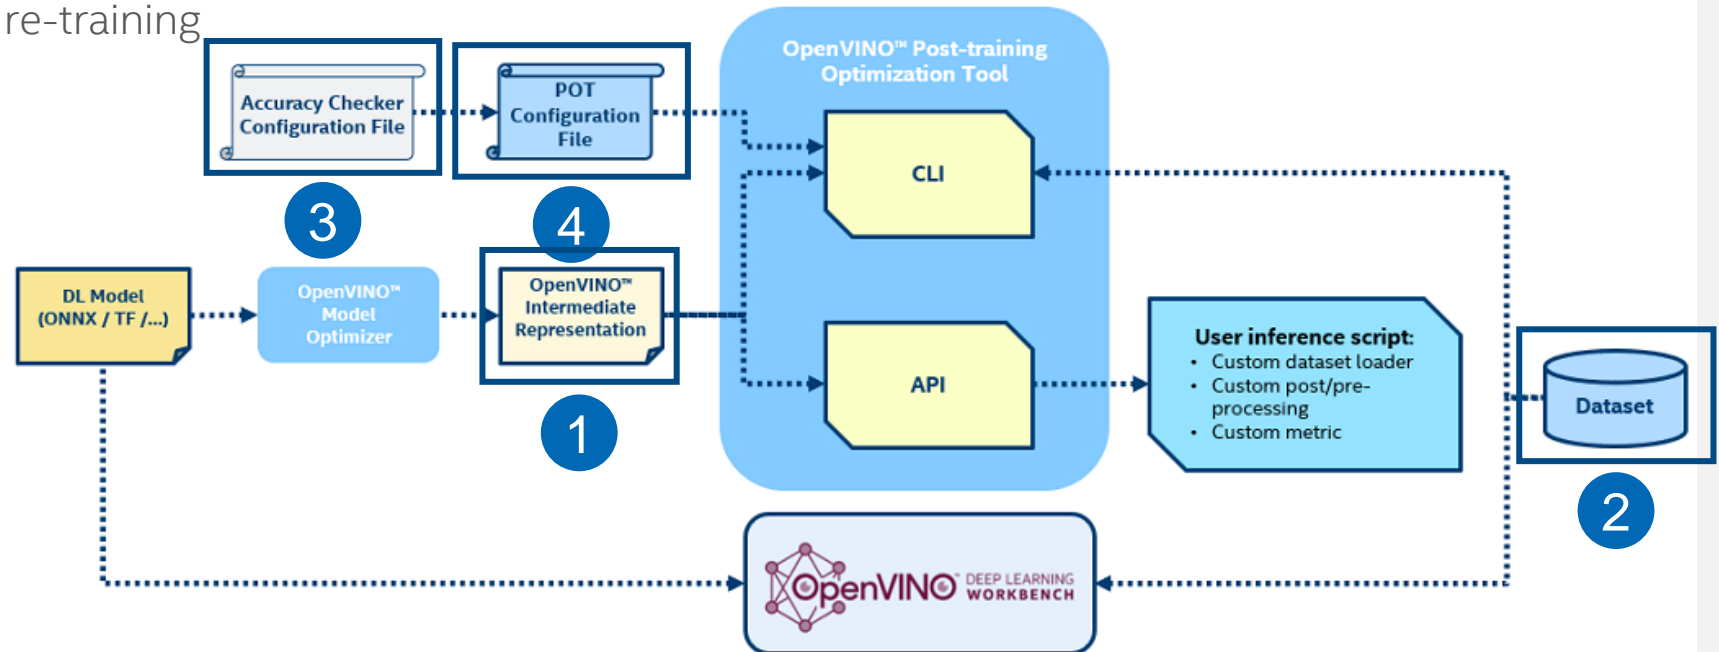
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

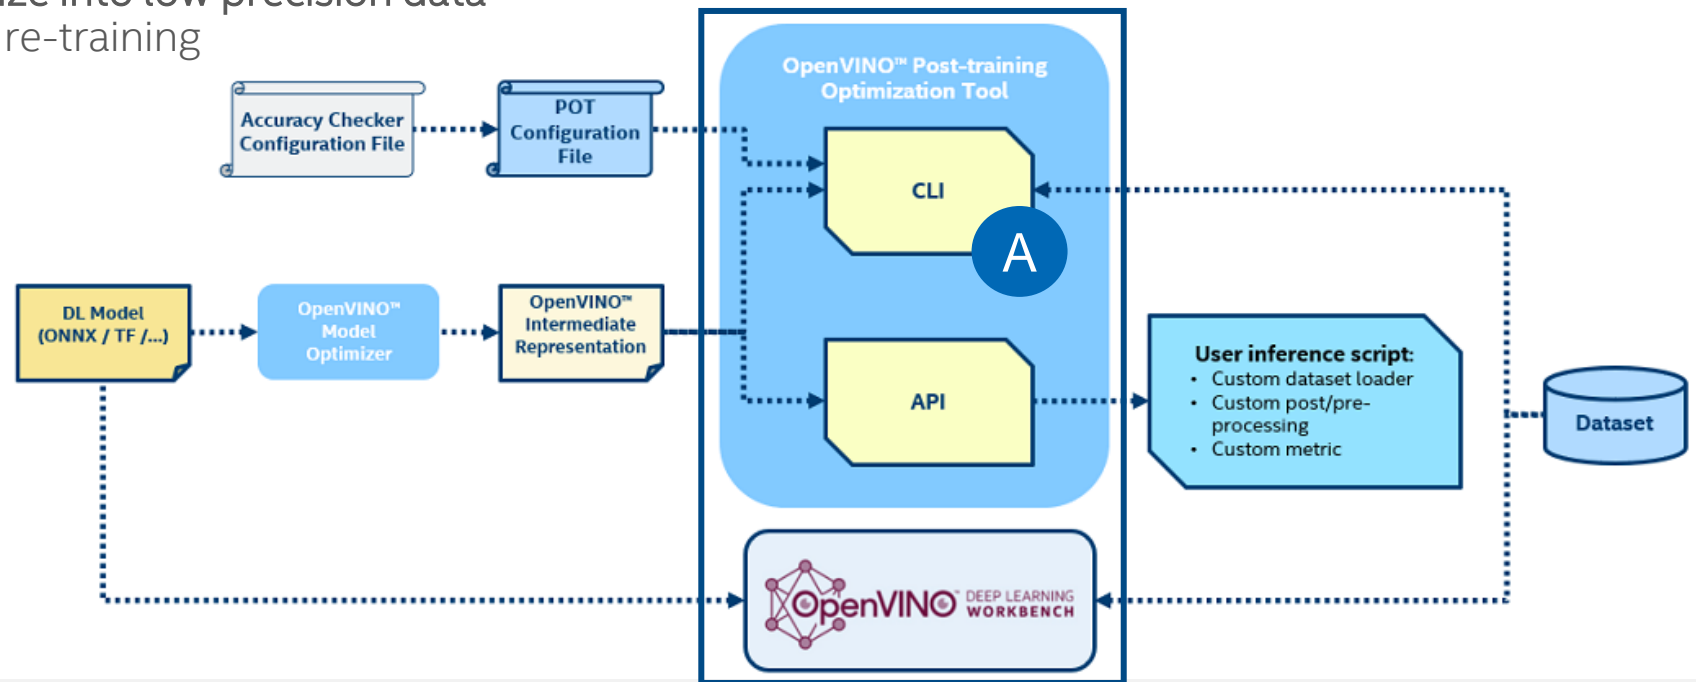
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

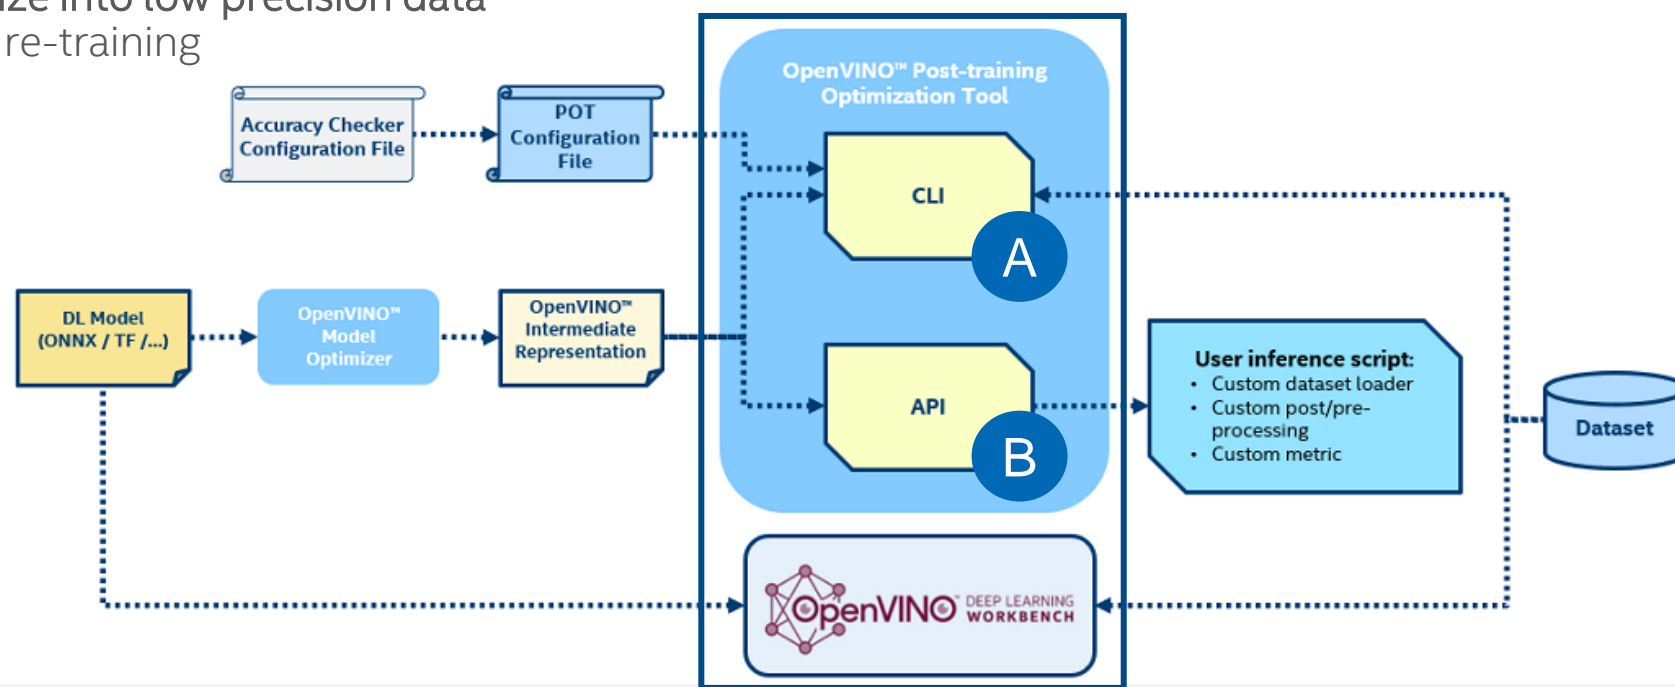
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

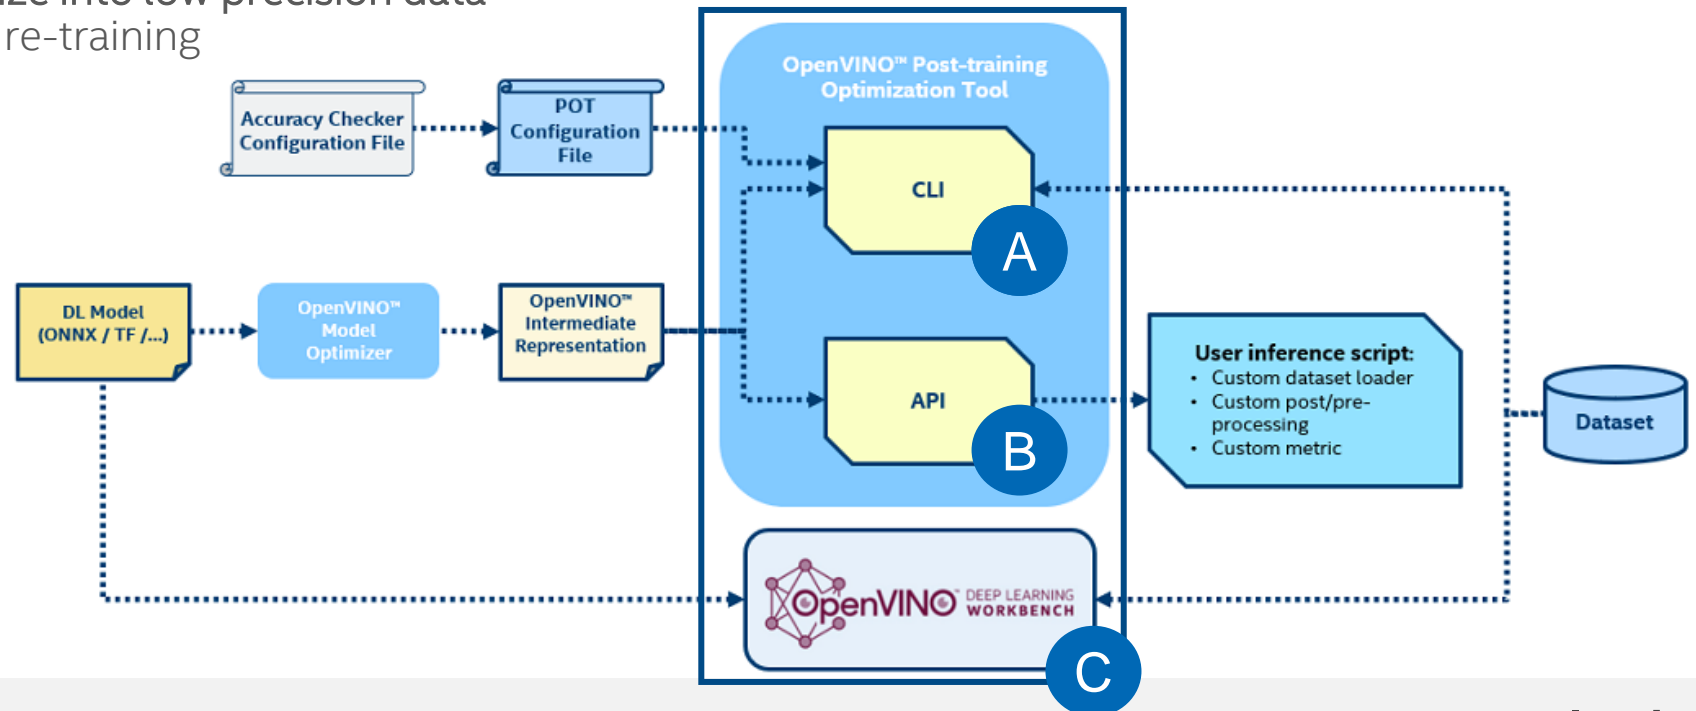
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

[https://docs.openvino toolkit.org/latest/pot\\_README.html](https://docs.openvino toolkit.org/latest/pot_README.html)

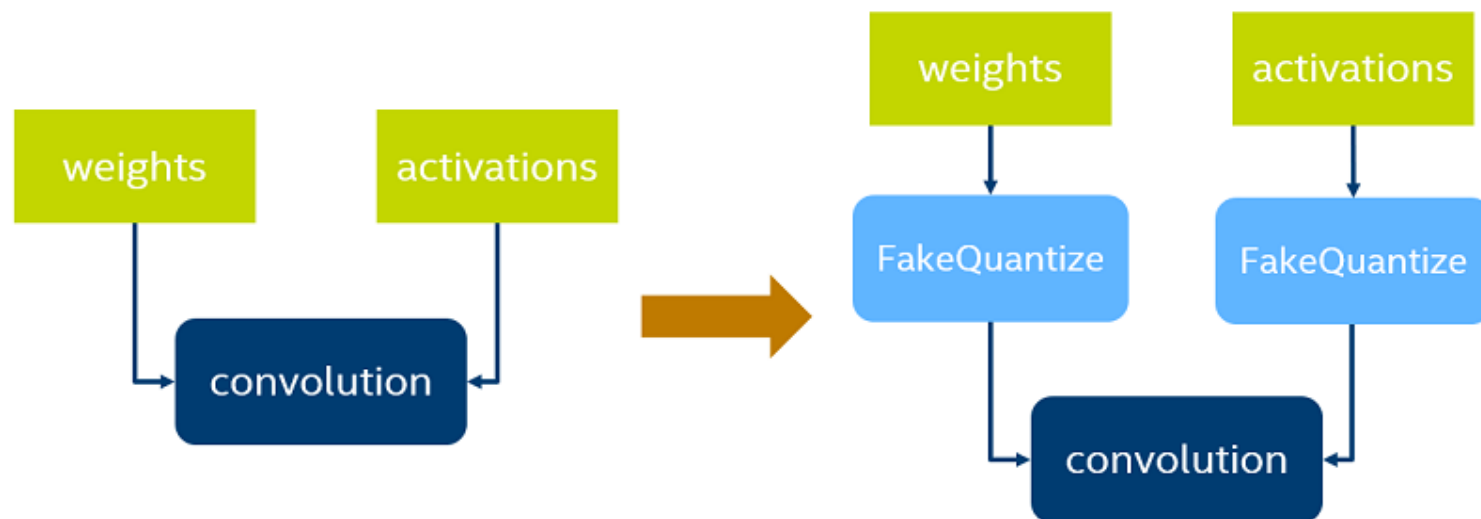
- Using the Python\* API, the Post-training Optimization Tool integrates with the Model Optimizer, DL Workbench and accuracy checker tools to streamline the development process
- Enables a conversion technique of deep learning model that **reduces model size into low precision data types**, such as INT8, without re-training
- Reduces model size **while also improving latency, with little degradation** in model accuracy and without model re-training.
- Different optimization approaches are supported: quantization algorithms, sparsity, etc.



# Post-Training Optimization Tool

## Quantization Algorithms

[https://docs.openvinotoolkit.org/latest/pot\\_compression\\_algorithms\\_quantization\\_README.html](https://docs.openvinotoolkit.org/latest/pot_compression_algorithms_quantization_README.html)




3. **Tree-Structured Parzen Estimator (TPE)** tries to provide best possible performance improvement. It requires even more time for quantization than **AccuracyAwareQuantization** but may lead to better performance improvement.

1. **DefaultQuantization** is a default method that provides fast and, in most cases, accurate results for 8-bit quantization.
2. **AccuracyAwareQuantization** enables remaining at a predefined range of accuracy drop after quantization at the cost of performance improvement. It may require more time for quantization.

# Post-Training Optimization Tool – Best Practices

Start with DefaultQuantization Algorithm with default settings



```
7{
8  /* Model parameters */
9
10  "model": {
11    "model_name": "model_name", // Model name
12    "model": "<MODEL_PATH>", // Path to model (.xml format)
13    "weights": "<PATH_TO_WEIGHTS>" // Path to weights (.bin format)
14  },
15
16  /* Parameters of the engine used for model inference */
17
18  "engine": {
19    "config": "<CONFIG_PATH>" // Path to Accuracy Checker config
20  },
21}
```

default\_quantization\_template.json

```
21
22  /* Optimization hyperparameters */
23
24  "compression": {
25    "target_device": "ANY", // Target device, the specificity of which will be taken
26                          // into account during optimization
27    "algorithms": [
28      {
29        "name": "DefaultQuantization", // Optimization algorithm name
30        "params": {
31          "preset": "performance", // Preset [performance, mixed, accuracy] which control the quantization
32                                // mode (symmetric, mixed (weights symmetric and activations asymmetric)
33                                // and fully asymmetric respectively)
34
35          "stat_subset_size": 300 // Size of subset to calculate activations statistics that can be used
36                                // for quantization parameters calculation
37        }
38      }
39    ]
40  }
41}
```

# Post-Training Optimization Tool – Best Practices

## Tuning Hyperparameters of the DefaultQuantization

- **"preset"** - preset which controls the quantization mode (symmetric and asymmetric). It can take two values:
  - **"performance"** (default) - stands for symmetric quantization of weights and activations. This is the most performant across all the HW.
  - **"mixed"** - symmetric quantization of weights and asymmetric quantization of activations. This mode can be useful for quantization of NN which has both negative and positive input values in quantizing operations, e.g. non-ReLU based CNN.
- **"stat\_subset\_size"** - size of subset to calculate activations statistics used for quantization. The whole dataset is used if no parameter specified. We recommend using not less than 300 samples.
- **"ignored"** - NN subgraphs which should be excluded from the optimization process
  - **"scope"** - list of particular nodes to exclude



# Post-Training Optimization Tool – Best Practices

## Tuning Hyperparameters of the DefaultQuantization

- `"preset"` - preset which controls the quantization mode (symmetric and asymmetric). It can take two values:
  - `"performance"` (default) - stands for symmetric quantization of weights and activations. This is the most performant across all the HW.
  - `"mixed"` - symmetric quantization of weights and asymmetric quantization of activations. This mode can be useful for quantization of NN which has both negative and positive input values in quantizing operations, e.g. non-ReLU based CNN.
- `"stat_subset_size"` - size of subset to calculate activations statistics used for quantization. The whole dataset is used if no parameter specified. We recommend using not less than 300 samples.
- `"ignored"` - NN subgraphs which should be excluded from the optimization process
  - `"scope"` - list of particular nodes to exclude

# Post-Training Optimization Tool – Best Practices

## Tuning Hyperparameters of the DefaultQuantization

- **"preset"** - preset which controls the quantization mode (symmetric and asymmetric). It can take two values:
  - **"performance"** (default) - stands for symmetric quantization of weights and activations. This is the most performant across all the HW.
  - **"mixed"** - symmetric quantization of weights and asymmetric quantization of activations. This mode can be useful for quantization of NN which has both negative and positive input values in quantizing operations, e.g. non-ReLU based CNN.
- **"stat\_subset\_size"** - size of subset to calculate activations statistics used for quantization. The whole dataset is used if no parameter specified. We recommend using not less than 300 samples.

- **"ignored"** - NN subgraphs which should be excluded from the optimization process
  - **"scope"** - list of particular nodes to exclude

# Post-Training Optimization Tool – Best Practices

## Tuning Hyperparameters of the DefaultQuantization

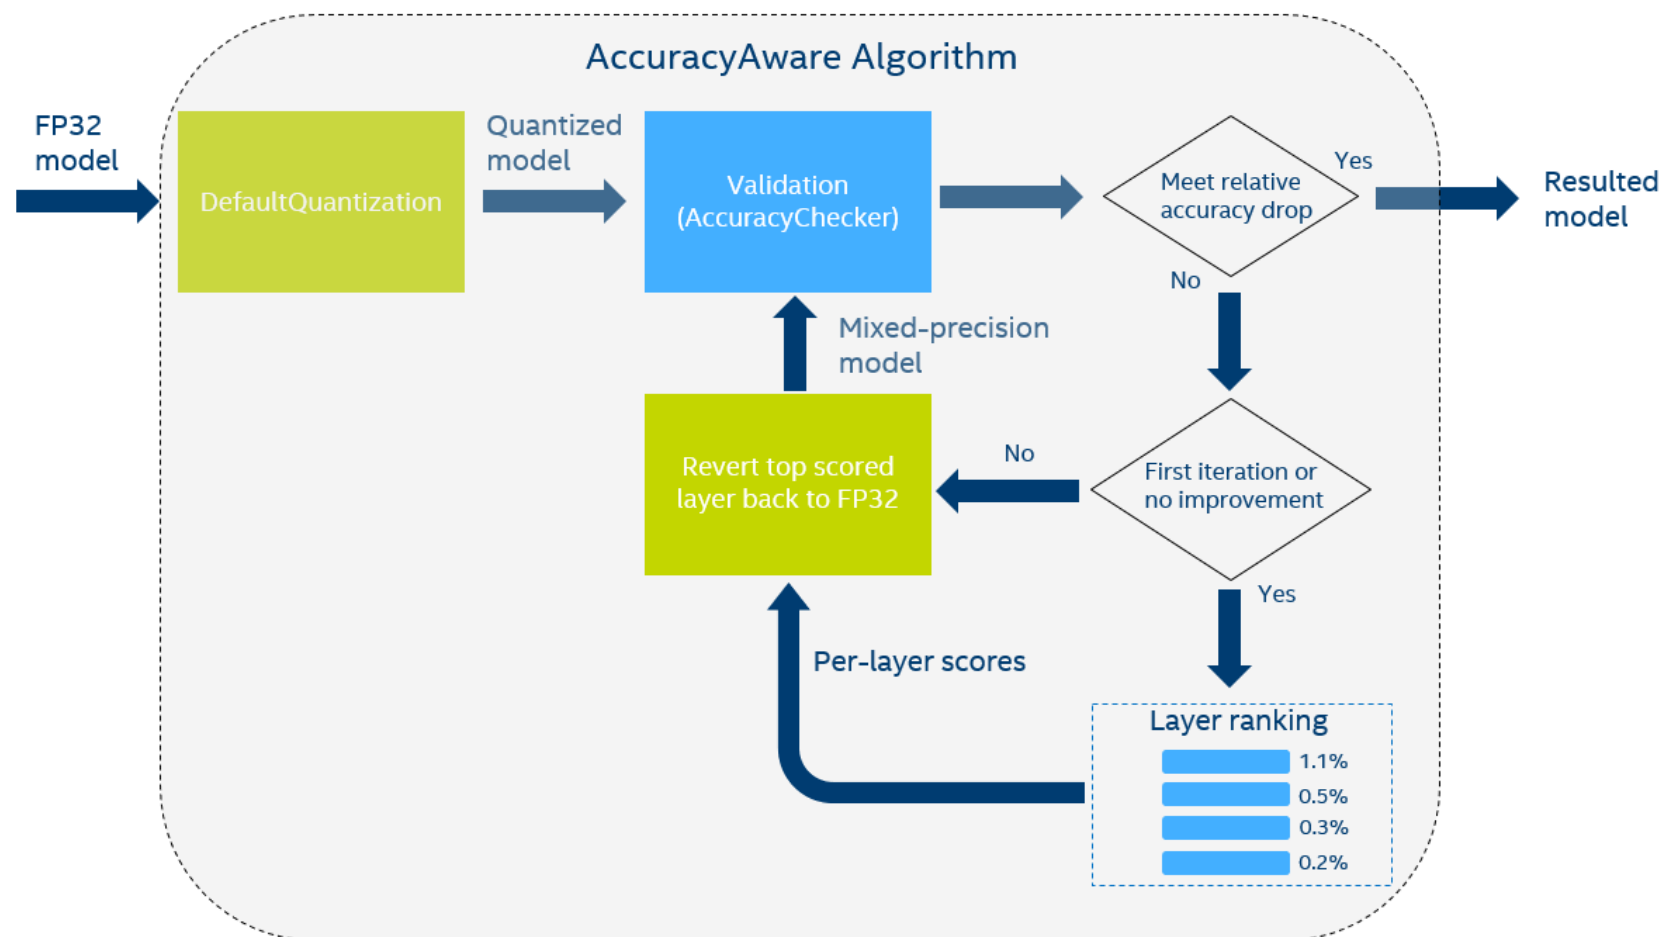
- `"range_estimator"` - this section describes parameters of range estimator that is used in MinMaxQuantization method to get the quantization ranges and filter outliers based on the collected statistics. These are the parameters that user can vary to get better accuracy results:
  - `"max"` - parameters to estimate top border of quantizing floating-point range:
    - `"type"` - type of the estimator:
      - `"max"` (default) - estimates the maximum in the quantizing set of value
      - `"quantile"` - estimates the quantile in the quantizing set of value
    - `"outlier_prob"` - outlier probability used in the "quantile" estimator
  - `"min"` - parameters to estimate bottom border of quantizing floating-point range:
    - `"type"` - type of the estimator:
      - `"min"` (default) - estimates the minimum in the quantizing set of value
      - `"quantile"` - estimates the quantile in the quantizing set of value
    - `"outlier_prob"` - outlier probability used in the "quantile" estimator

**Learn more about  
DefaultQuantization:**

[https://docs.openvinotoolkit.org/latest/pot\\_compression\\_algorithms\\_quantization\\_default\\_README.html](https://docs.openvinotoolkit.org/latest/pot_compression_algorithms_quantization_default_README.html)

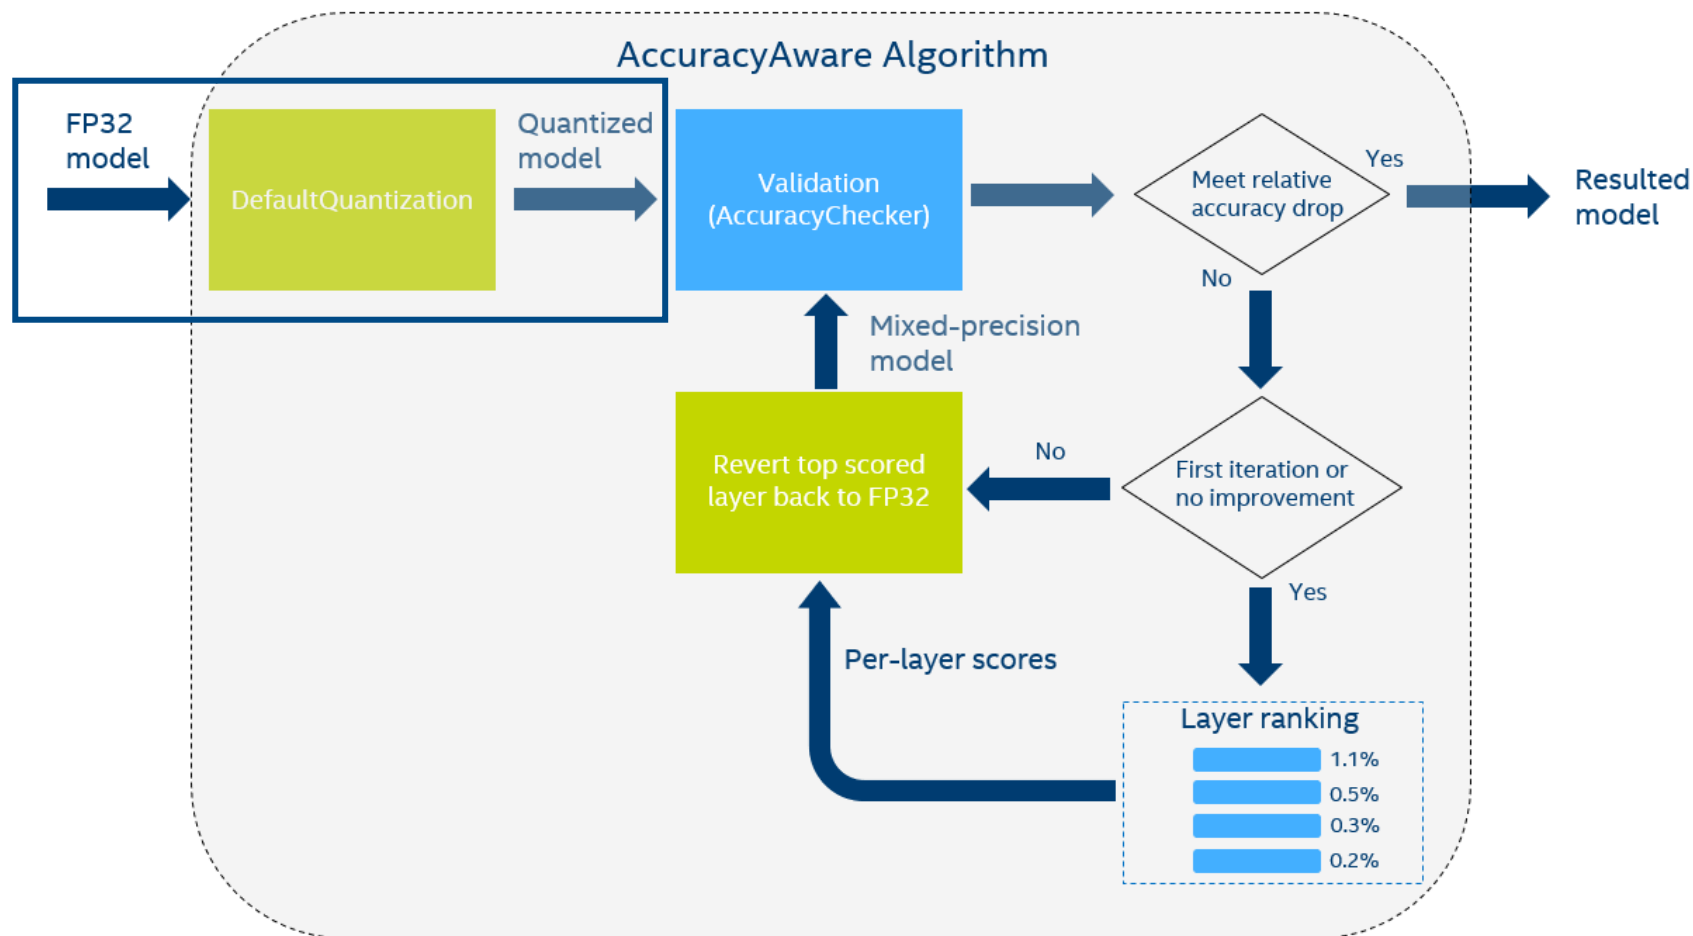
# Post-Training Optimization Tool – Best Practices

Try AccuracyAwareQuantization Method if not satisfied with previous steps



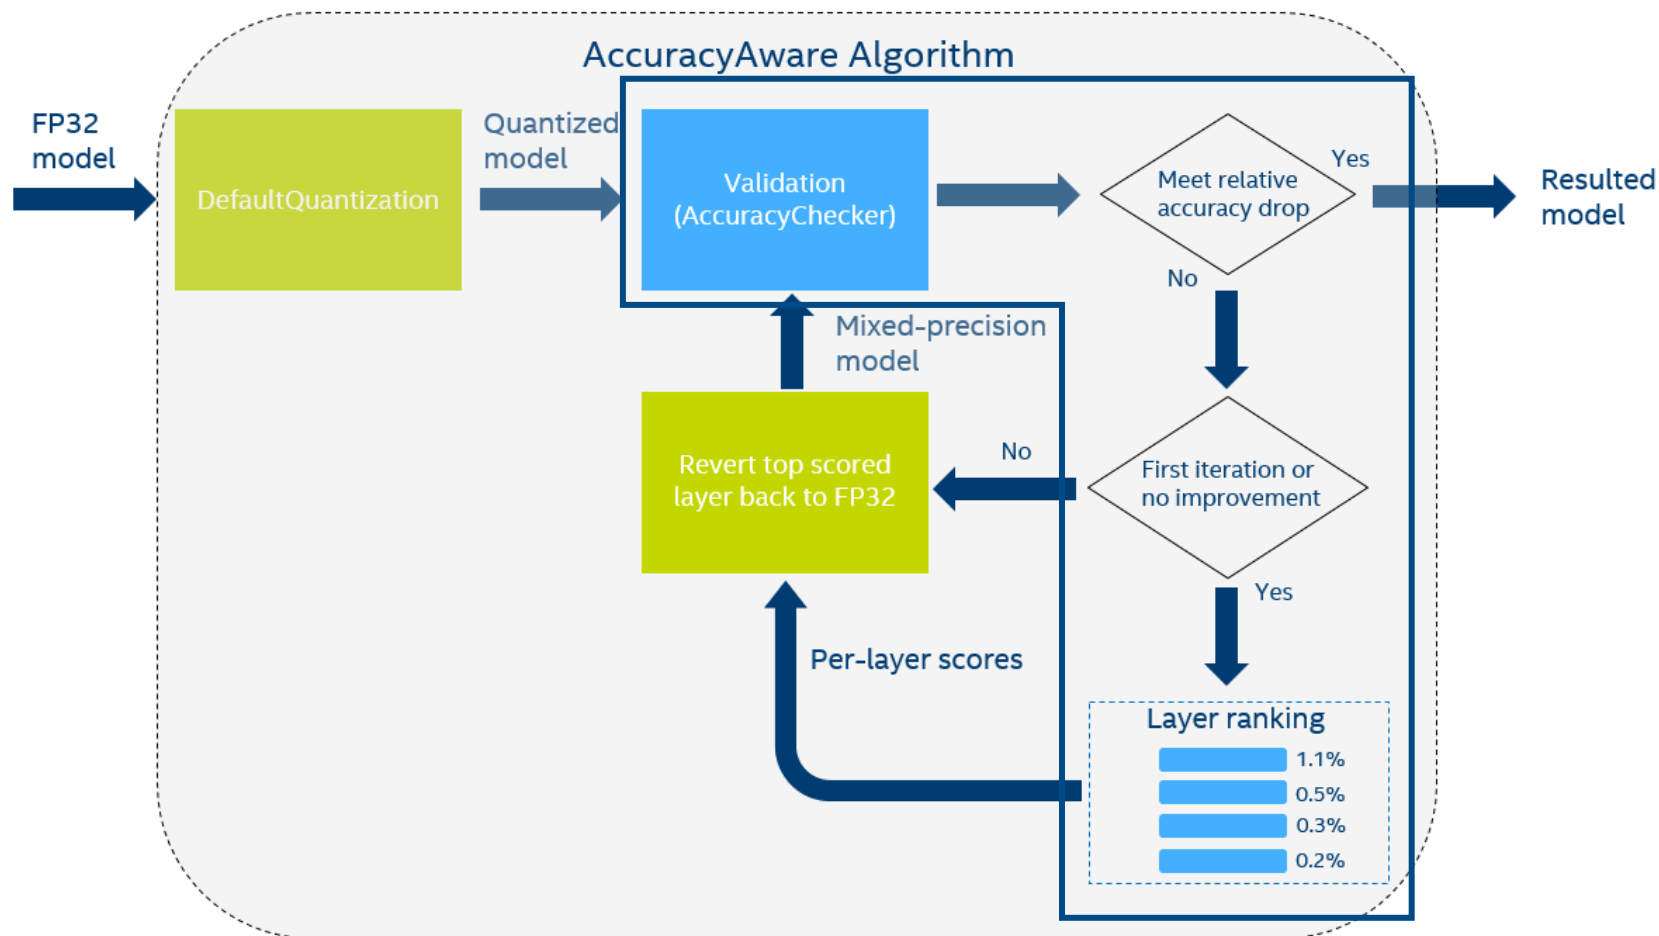
# Post-Training Optimization Tool – Best Practices

Try AccuracyAwareQuantization Method if not satisfied with previous steps



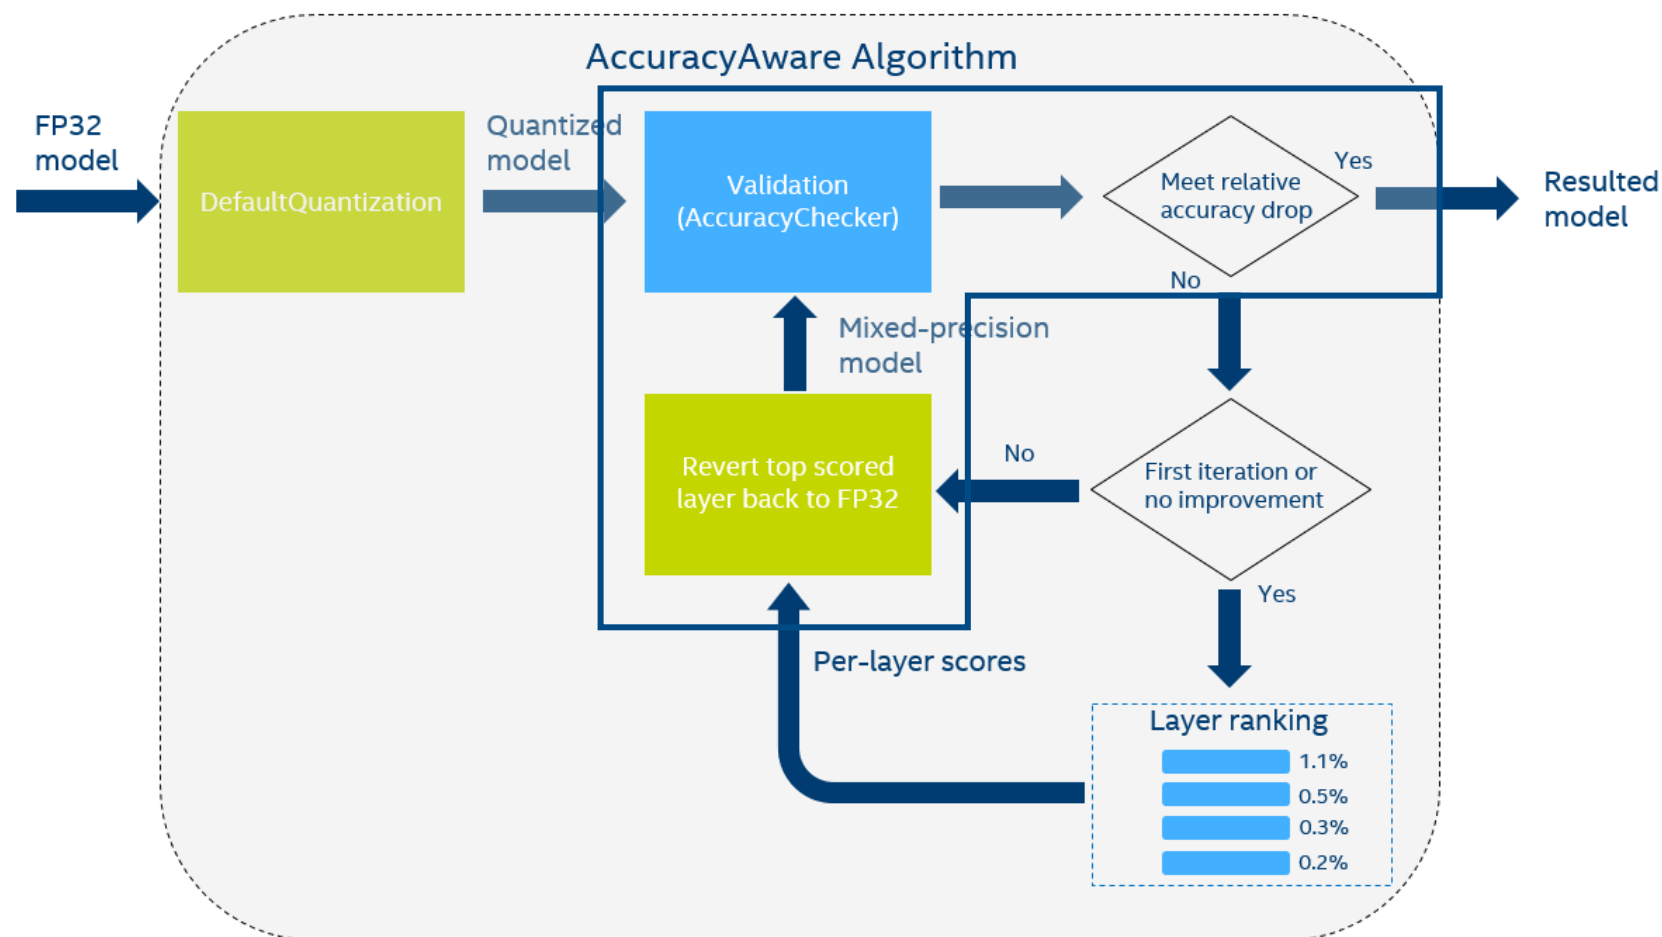
# Post-Training Optimization Tool – Best Practices

Try AccuracyAwareQuantization Method if not satisfied with previous steps



# Post-Training Optimization Tool – Best Practices

Try AccuracyAwareQuantization Method if not satisfied with previous steps



# Post-Training Optimization Tool – Best Practices

Try AccuracyAwareQuantization Method if not satisfied with previous steps

```
6
7 {
8   /* Model parameters */
9
10  "model": {
11    "model_name": "model_name", // Model name
12    "model": "<MODEL_PATH>", // Path to model (.xml format)
13    "weights": "<PATH_TO_WEIGHTS>" // Path to weights (.bin format)
14  },
15
16  /* Parameters of the engine used for model inference */
17
18  "engine": {
19    "config": "<CONFIG_PATH>" // Path to Accuracy Checker config
20  },
21
22  /* Optimization hyperparameters */
23
24  "compression": {
25    "target_device": "ANY", // Target device, the specificity of which will be taken
26                          // into account during optimization
27    "algorithms": [
28      {
29        "name": "AccuracyAwareQuantization", // Optimization algorithm name
30        "params": {
31          "preset": "performance", // Preset [performance, mixed, accuracy] which control the quantization
32                                // mode (symmetric, mixed (weights symmetric and activations asymmetric)
33                                // and fully asymmetric respectively)
34
35          "stat_subset_size": 300, // Size of subset to calculate activations statistics that can be used
36                                // for quantization parameters calculation
37
38          "maximal_drop": 0.01, // Maximum accuracy drop which has to be achieved after the quantization
39
40          "tune_hyperparams": false // Whether to search the best quantization parameters for model
41        }
42      }
43    ]
44  }
45 }
```

accuracy\_aware\_quantization\_template.json

Learn more about  
AccuracyAwareQuantization :

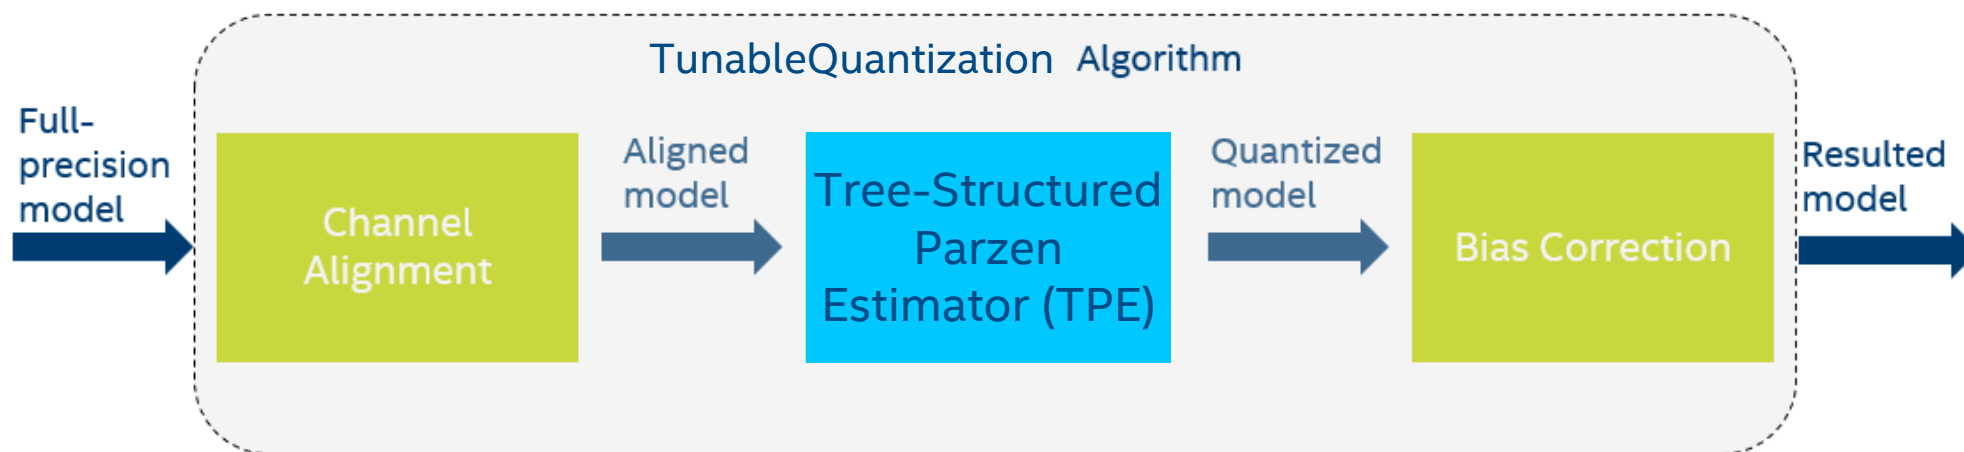
[https://docs.openvinotoolkit.org/latest/pot\\_compression\\_algorithms\\_quantization\\_accuracy\\_aware\\_README.html](https://docs.openvinotoolkit.org/latest/pot_compression_algorithms_quantization_accuracy_aware_README.html)



# Post-Training Optimization Tool – Best Practices

Lastly: **TunableQuantization** with layer-Wise Hyperparameters Tuning Using TPE (Tree of Parzen Estimators )

- TunableQuantization algorithm is a modified version (to support hyperparameters setting by **Tree-Structured Parzen Estimator (TPE)**) of the vanilla MinMaxQuantization quantization method that automatically inserts FakeQuantize operations into the model graph based on the specified target hardware and initializes them using statistics collected on the calibration dataset.



# Post-Training Optimization Tool – Best Practices

Lastly: TunableQuantization with layer-Wise Hyperparameters Tuning Using TPE (Tree of Parzen Estimators)

```
10 "model": {
11     "model_name": "model_name", // Model name
12     "model": "<MODEL_PATH>", // Path to a model (.xml format)
13     "weights": "<PATH_TO_WEIGHTS>" // Path to weights (.bin format)
14 },
15
16 /* Parameters of the engine used for model inference. */
17
18 "engine": {
19     "config": "<CONFIG_PATH>" // Path to Accuracy Checker config
20 },
21
22 /* Optimizer used to find "optimal" hyperparameters */
23
24 "optimizer": {
25     "name": "Tpe", // Global optimizer name
26     "params": {
27         "max_trials": 200, // Maximum number of trials
28         "trials_load_method": "cold start", // Start from scratch or
29         "accuracy_loss": 0.1, // Accuracy threshold (%)
30         "latency_reduce": 1.5, // Target latency improvement versus
31         "accuracy_weight": 1.0, // Accuracy weight in loss function
32         "latency_weight": 1.0, // Latency weight in loss function
33         "benchmark": {
34             // Latency measurement benchmark configuration (https://
35             "performance_count": false,
36             "batch_size": 0,
37             "nthreads": 4,
38             "nstreams": 0,
39             "nreq": 0,
40             "api_type": "sync",
41             "niter": 4,
42             "duration_seconds": 30,
43             "benchmark_app_dir": "<path to benchmark app>" // Path to
44         }
45     }
46 }
47
48 "compression": {
49     "target_device": "ANY", // Target device, the spec
50                             // into account during opt
51     "algorithms": [
52         {
53             "name": "ActivationChannelAlignment",
54             "params": {
55                 "stat_subset_size": 300 // Size of sub
56                                         // for quantiz
57             }
58         },
59         {
60             "name": "TunableQuantization",
61             "params": {
62                 /* Preset is a collection of optimizat
63                 to improve which metric the algorithm
64                 [performance, mixed, accuracy] presets
65                 (symmetric, mixed(weights symmetric an
66                 "preset": "performance",
67                 "stat_subset_size": 300, // Size of
68                                         // for quan
69                 "tuning_scope": ["layer"] // List of c
70                                         // available
71             }
72         },
73         {
74             "name": "FastBiasCorrection",
75             "params": {
76                 "stat_subset_size": 300 // Size of sub
77                                         // for quantiz
78             }
79         }
80     ]
81 }
82
83 }
```

Learn more about  
TunableQuantization :

[https://docs.openvinotoolkit.org/latest/pot\\_compression\\_algorithms\\_quantization\\_tunable\\_quantization\\_README.html](https://docs.openvinotoolkit.org/latest/pot_compression_algorithms_quantization_tunable_quantization_README.html)

Learn more about Tree-  
Structured Parzen  
Estimator (TPE) :

[https://docs.openvinotoolkit.org/latest/pot\\_compression\\_optimization\\_tpe\\_README.html](https://docs.openvinotoolkit.org/latest/pot_compression_optimization_tpe_README.html)

# Deep Learning Workbench

July 2021



intel<sup>®</sup>

# Deep Learning Workbench

[https://docs.openvino toolkit.org/latest/workbench\\_docs\\_Workbench\\_DG\\_Introduction.html](https://docs.openvino toolkit.org/latest/workbench_docs_Workbench_DG_Introduction.html)

- Web-based, UI extension tool of the Intel® Distribution of OpenVINO™ toolkit
- Visualizes performance data for topologies and layers to aid in model analysis
- Automates analysis for optimal performance configuration (streams, batches, latency)
- Experiment with INT8 or Winograd calibration for optimal tuning using the Post Training Optimization Tool
- Provide accuracy information through accuracy checker
- Direct access to models from public set of Open Model Zoo
- Enables remote profiling, allowing the collection of performance data from multiple different machines without any additional set-up.

## Development Guide ▶

[https://docs.openvino toolkit.org/latest/docs\\_Workbench\\_DG\\_Introduction.html](https://docs.openvino toolkit.org/latest/docs_Workbench_DG_Introduction.html)

### Work with models in Deep Learning Workbench

Obtain for specific use case

Explore inference configurations

Optimize for higher performance

Measure accuracy

Visualize model output

Prepare for deployment

### Learn OpenVINO in Deep Learning Workbench

Analyze performance

Explore model architecture

Experiment with OpenVINO API

Follow full OpenVINO workflow created for your model

# Installation Methods

## Run the Deep Learning Workbench Locally

This section contains instructions on how to run DL Workbench. Select your options and run the commands on your local machine. Please, ensure that you have met the [prerequisites](#).

### General Options

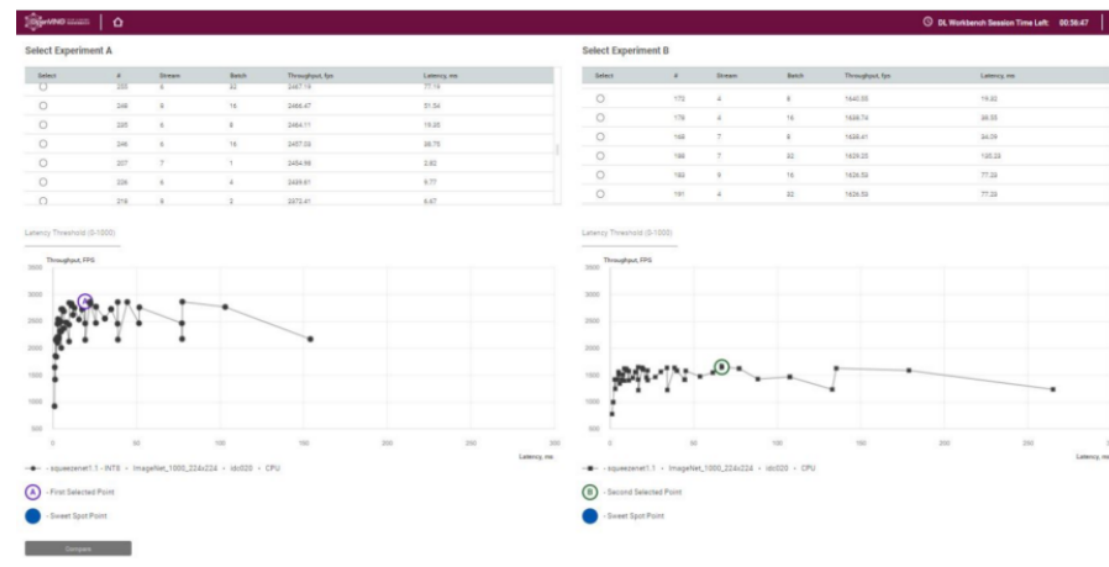
I have Docker installed	<input type="radio"/> Yes	<input checked="" type="radio"/> No	
My OS	<input checked="" type="radio"/> Linux	<input type="radio"/> Windows	<input type="radio"/> macOS
Accelerators on my machine	<input checked="" type="checkbox"/> CPU	<input checked="" type="checkbox"/> GPU	<input checked="" type="checkbox"/> NCS2 <input type="checkbox"/> HDDL

### Advanced Options

Start DL Workbench with	<input checked="" type="radio"/> Python wrapper	<input type="radio"/> Docker command	
Use HTTP Proxy	<input type="radio"/> Yes	<input checked="" type="radio"/> No	Enter your proxy
Use HTTPS Proxy	<input type="radio"/> Yes	<input checked="" type="radio"/> No	Enter your proxy
Use No Proxy	<input type="radio"/> Yes	<input checked="" type="radio"/> No	Enter your proxy

### Execute / Results

1. Install Docker	Follow <a href="#">instructions</a> on configuring Docker on Linux
2. Install Python Wrapper	<code>python3 -m pip install -U openvino-workbench</code>
3. Start DL Workbench	<code>openvino-workbench --image openvino/workbench:2021.4 --enable-gpu --enable-myriad</code>



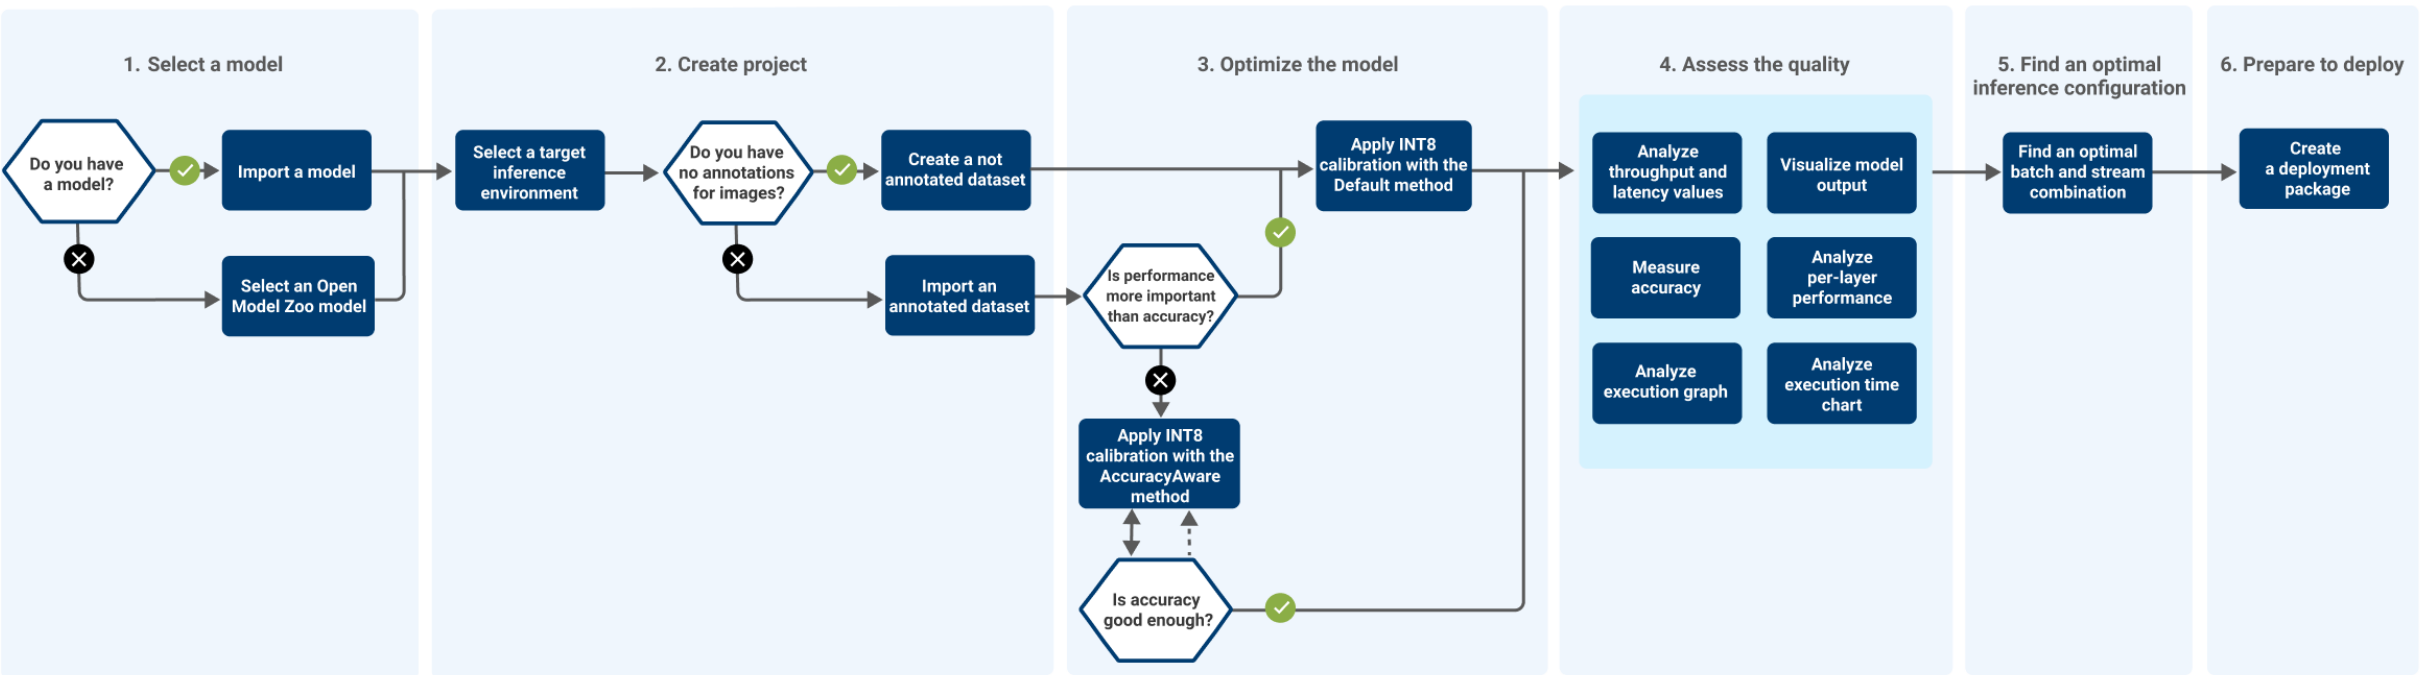
## Deep Learning Workbench

The Deep Learning Workbench simplifies using the Intel® Distribution of OpenVINO toolkit to tune, visualize, and compare the performance of deep learning models on Intel® architecture.

[Go There](#) →

**Note:** To get full features of DL Workbench, please run it on local system

# Deep Learning Workbench Workflow



# DL Workbench Demo

# Deep Learning Streamer

April 2021



intel®




# Introducing.. DL streamer

- Intel® Distribution of OpenVINO™ toolkit [Deep Learning \(DL\) Streamer](#), now part of the default installation package
- Enables developers to **create and deploy** optimized streaming media analytics **pipelines** across Intel® architecture from edge to cloud
- Optimal pipeline interoperability with a **familiar developer experience** built using the GStreamer multimedia framework



# What is GStreamer?

- A pipeline consists of **connected processing elements**
- Each element is provided by a **plug-in** and can be **grouped into bins**
- Elements communicate by means of **pads** – source pad and sink pad
- Data buffers flow **from Source element to Sink element** & from source pad to sink pad



Ref:  
<https://gstreamer.freedesktop.org/data/doc/gstreamer/head/manual/manual.pdf>

# Under the hood: DL Streamer

Application

Reference Application Designs

GStreamer framework

GStreamer  
plugins

GStreamer Media Plugins (Standard)

Decode

VPP

Encode

DL Streamer - GStreamer Video Analytics (GVA)  
Plugin

Detect

Classify

Track

Publish

Runtime  
Libraries

VAAPI

Libav

Intel® Distribution of OpenVINO™  
toolkit Deep Learning  
Inference Engine

OpenCV

MQTT/  
Kafka

Hardware



# Media Processing Pipeline

Video Pipeline – decode, convert, render

filesrc — decodebin — videoconvert — xvimagesink

input

HW/SW  
decode

convert


render  
on screen



```
gst-launch-1.0 filesrc location=/path/to/video.mp4 ! decodebin ! videoconvert ! xvimagesink
```

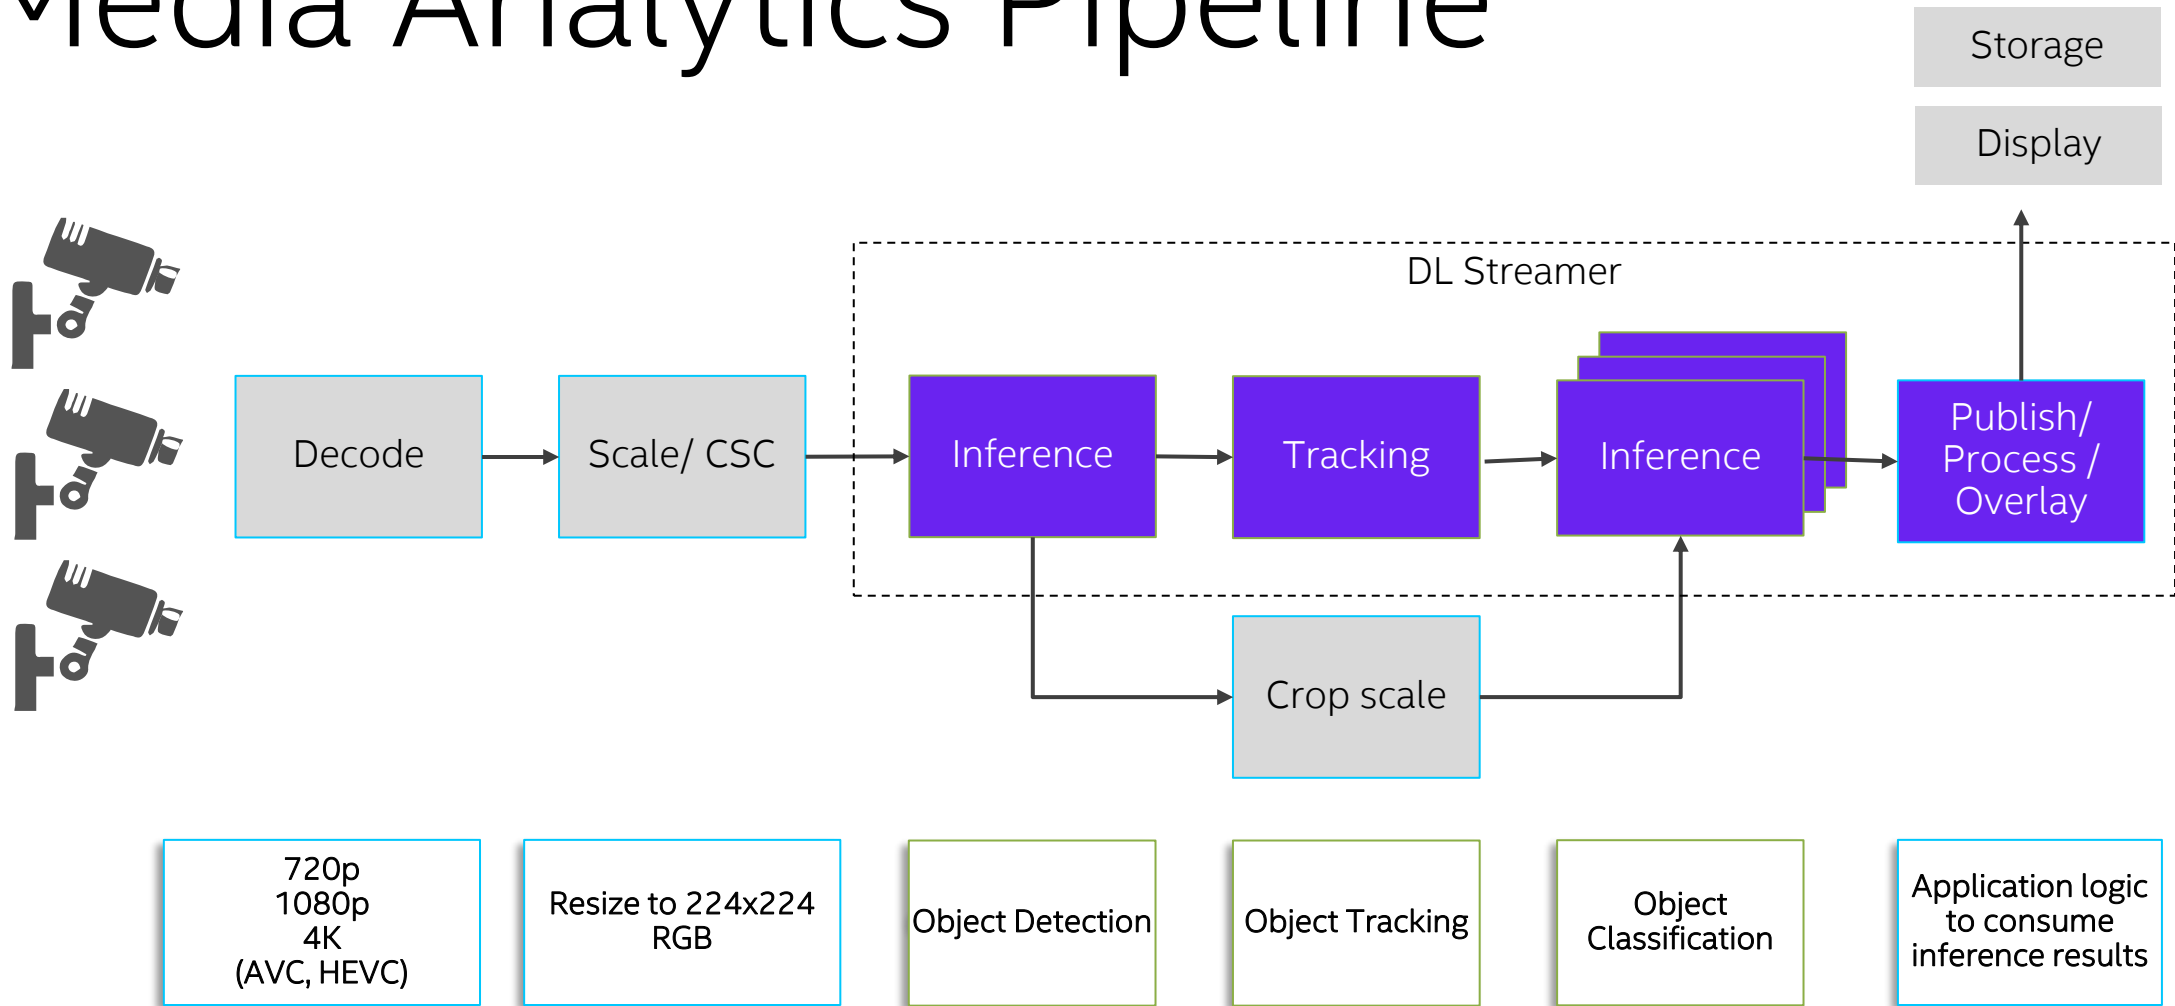
# Using the DL Streamer

Video Analytics pipeline – person and vehicle detection, person, vehicle attributes classification

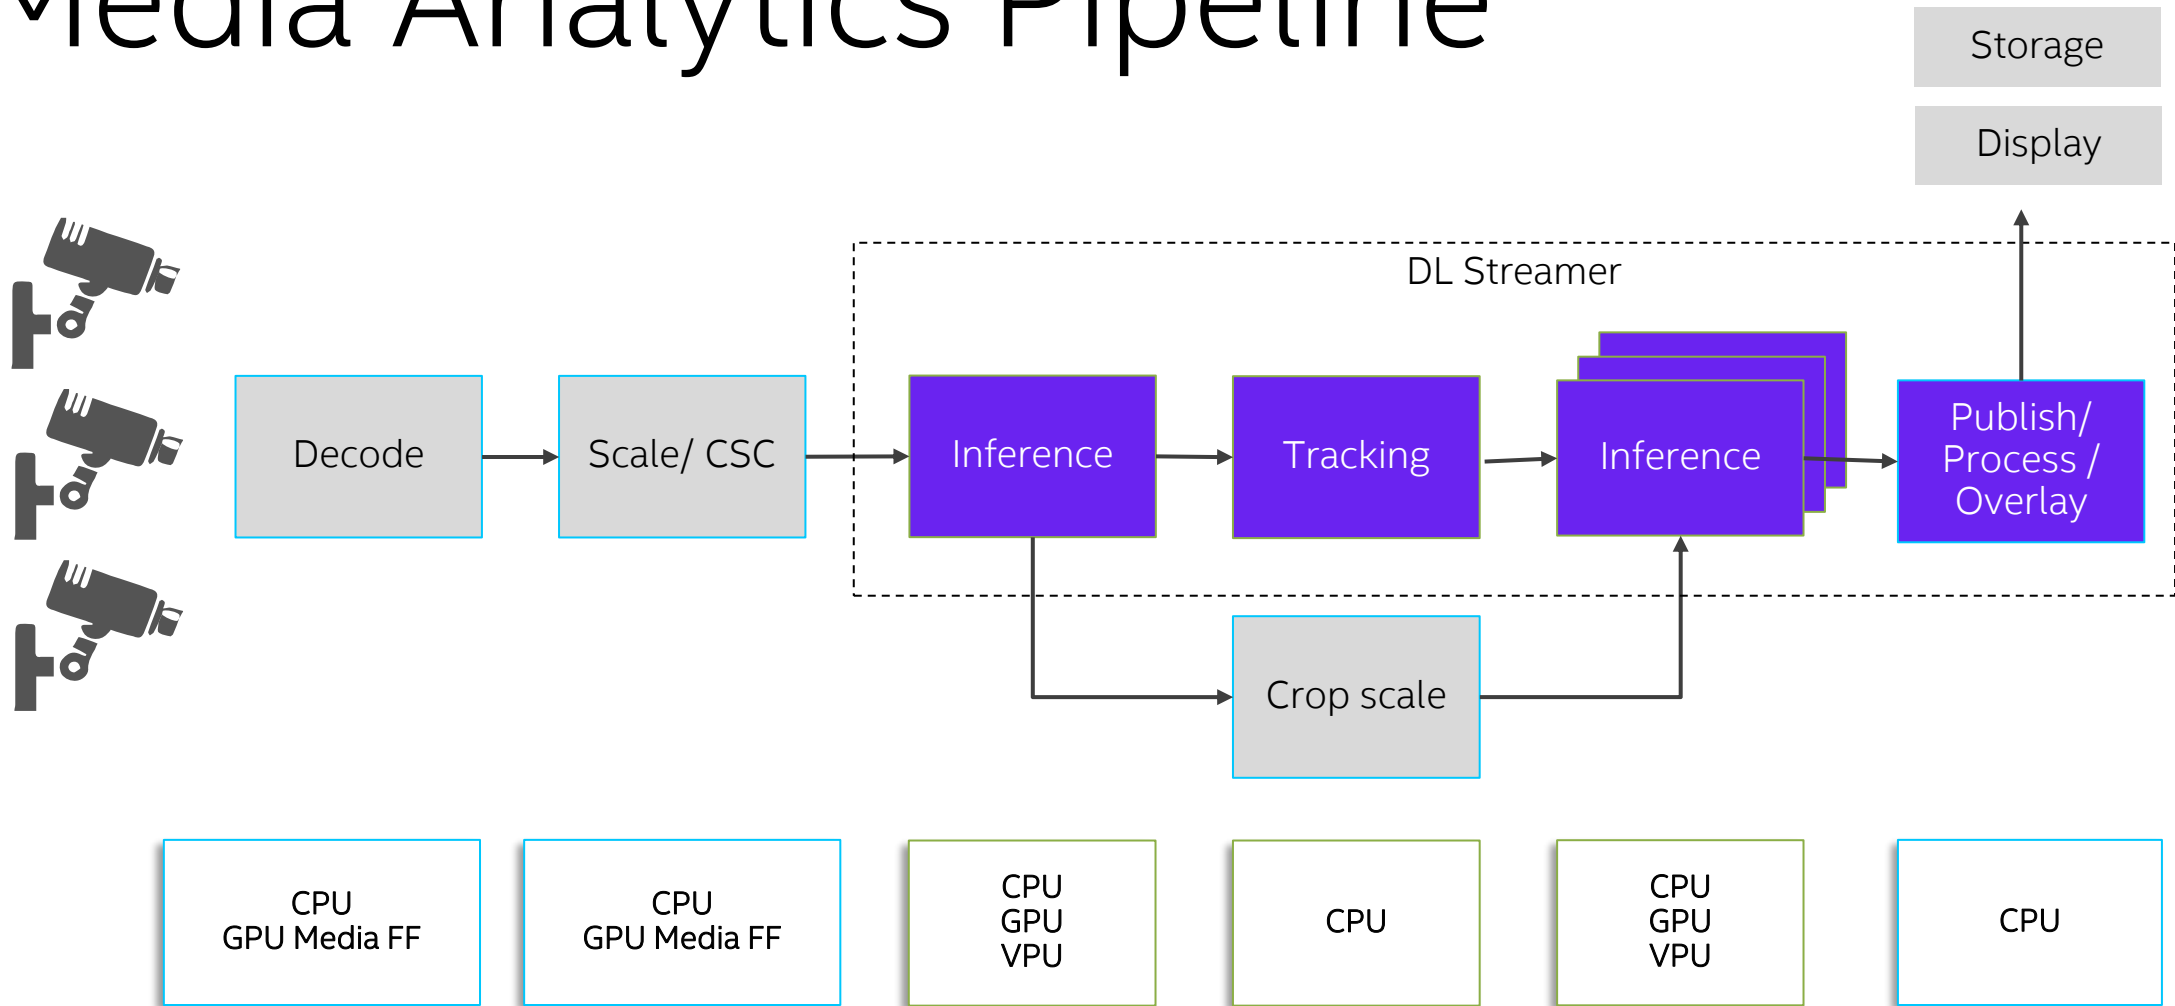


```
gst-launch-1.0 filesrc location=/path/to/video.mp4 !
decodebin ! videoconvert ! video/x-raw,format=BGRx ! \
gvadetect model=person-vehicle-bike-detection-crossroad-0078.xml model-proc=person-vehicle-bike-detection-
crossroad-0078.json inference-interval=10 threshold=0.6 device=CPU ! queue ! \
gvatrack tracking-type="short-term" ! queue ! \
gvaclassify model= person-attributes-recognition-crossroad-0230.xml model-proc= person-attributes-recognition-
crossroad-0230.json reclassify-interval=10 device=CPU object-class=person ! queue ! \
gvaclassify model= vehicle-attributes-recognition-barrier-0039.xml model-proc= vehicle-attributes-recognition-
barrier-0039.json reclassify-interval=10 device=CPU object-class=vehicle ! queue ! \
gvawatermark ! videoconvert ! fpsdisplaysink video-sink=xvimagesink sync=true
```

# Media Analytics Pipeline



# Media Analytics Pipeline



# Audio Processing

## DL Streamer for end-to-end audio analytics pipeline



- Intel® Distribution of OpenVINO™ toolkit [Deep Learning \(DL\) Streamer](#), part of the default installation package
- Enables developers to create and deploy optimized streaming media analytics pipelines across Intel® architecture from edge to cloud
- Optimal pipeline interoperability with a familiar developer experience built using the GStreamer\* multimedia framework
- Introduces `gva audiodetect` for audio event detection
  - Can be paired with `alcnet` public model for end-to-end audio analytics pipeline

### DL Streamer Elements:

- [gva audiodetect](#) for audio event detection using ACLNet
- [gvametaconvert](#) for converting ACLNet detection results into JSON for further processing and display
- [gvametapublish](#) for printing detection results to stdout



# Resources to Get Started



Intel® Distribution of OpenVINO™ Toolkit:

<https://software.intel.com/content/www/us/en/develop/tools/opencvino-toolkit.html>

Intel® Edge Software Hub

Download prevalidated software to learn, develop, and test your solutions for the edge.

Intel® Edge Software Hub:

<https://software.intel.com/content/www/us/en/develop/topics/iot/edge-solutions.html>

Intel® DevCloud  
FOR THE EDGE

Intel® DevCloud for the Edge:

<https://devcloud.intel.com/edge/home>

To get access to the full video series, please complete the short form: <http://intel.ly/38B9ix6>

